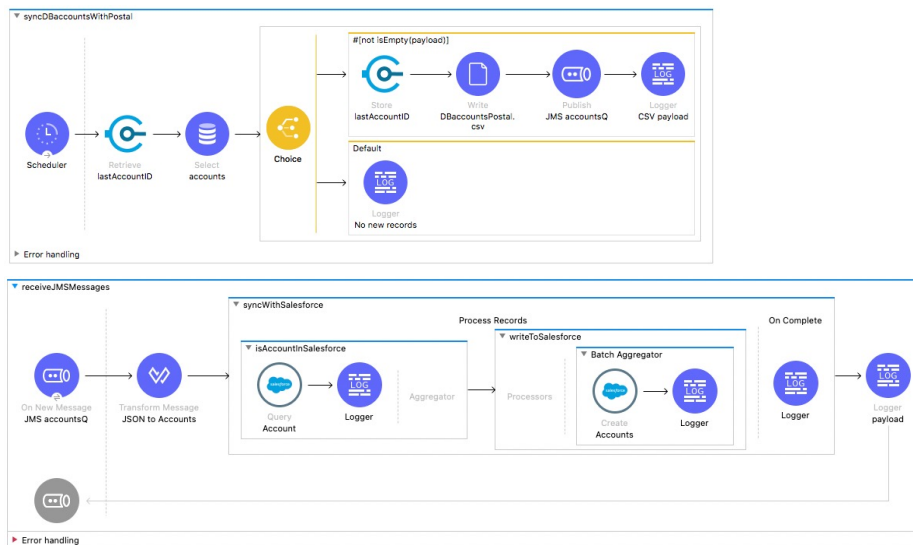




PART 3: Building Applications to Synchronize Data

1


Goal



2

2

At the end of this part, you should be able to



- Trigger flows when files or database records are added or updated
- Schedule flows
- Persist and share data across flow executions
- Publish and consume JMS messages
- Process items in a collection sequentially
- Process records asynchronously in batches

All contents © MuleSoft Inc.

3

3



Module 12: Triggering flows




4

Goal



How have we initiated flows so far?

In this module, we will learn new ways




All contents © MuleSoft Inc.

5

5

At the end of this module, you should be able to



- Read and write files
- Trigger flows when files are added, created, or updated
- Trigger flows when new records are added to a database table
- Schedule flows to run at a certain time or frequency
- Persist and share data in flows using the Object Store
- Publish and consume JMS messages

All contents © MuleSoft Inc.

6

6

Reading and writing files

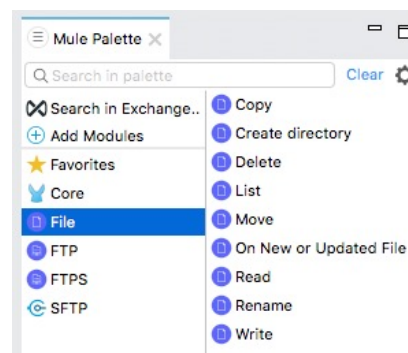


7

Reading and writing files



- There are 4 connectors for working with files and folders
 - File (for locally mounted file system)
 - FTP
 - FTPS
 - SFTP
- All have the same set of operations and they behave almost identically
- Support for
 - File matching functionality
 - Locking files
 - Overwriting, appending, and generating new files



All contents © MuleSoft Inc.


8

8

Using the File connector



- Add the File module to the project
- Create a global element configuration
 - Not required but a best practice
 - Set the **working directory** that will be the root for every path used with the connector
- Use one of the connector operations and specify its properties
- On **CloudHub**, the connector can only be used with the /tmp folder
- On **Customer-hosted** Mule runtimes, the account running Mule must have read and/or write permissions on the specified directories
- *Be careful not to permanently delete or overwrite files*
 - Move or rename them after processing




All contents © MuleSoft Inc.

9

9

Trigger a flow when a new file is created or updated



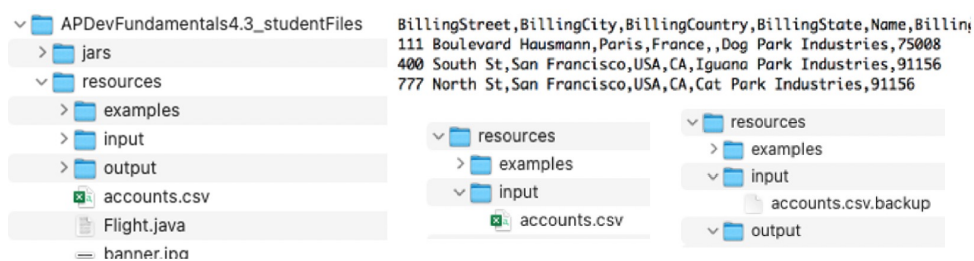
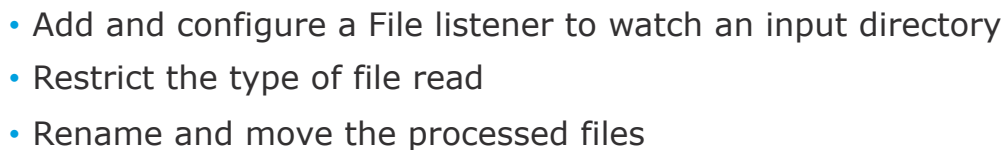
- Use the **On New or Updated File** listener
 - Polls a directory for files that have been created or updated
 - One message is generated for each file that is found
- Multiple ways to ensure a file is new
 - Delete each file after it has been processed so all files in the next poll will be new
 - Move each file to a different directory after it has been processed
 - Rename a file after it has been processed and filter the files to be processed
 - Save and compare the file creation or modification times



All contents © MuleSoft Inc.

10

10



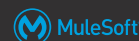
11

Synchronizing data with watermarks



6

Synchronizing data from one system to another



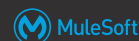
- The general process
 - The first time, you need to sync all the data
 - After that, you only need to sync the new data
 - Requires a field with ordered values to identify processed items
- How do you determine what is new and needs to be synced?
 - On the first sync, store the highest field value for any item in the data set
 - On later syncs, retrieve that value and compare the value of each item and see if it is larger
- The field with ordered values is often a
 - Record ID
 - Creation or modification timestamp

All contents © MuleSoft Inc.

13

13

Introducing watermarks



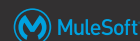
- The timestamp or ID that is stored each sync and then retrieved and compared against in the next sync is called a **watermark**
- Where did the name come from?
 - After a flood, one might record how high the water got by marking the level on a wall
 - Similarly, for data, we want to look at the last value - how “high” it was in the last sync

All contents © MuleSoft Inc.

14

14

Types of watermarking in Mule



- **Automatic**

- The saving, retrieving, and comparing is automatically handled for you
- Available for several connector listeners
 - On New or Updated File
 - On Table Row
- Restricted in how you can specify what items/records are retrieved

- **Manual**

- You handle saving, retrieving, and comparing the watermark
- More flexible in that you specify exactly what records you want retrieved

All contents © MuleSoft Inc.

15


15

Using listeners with automatic watermarking



16

Using automatic watermarking with files




- There is a watermarking option for the **On New and Updated File** operation for the family of file connectors



- There are two watermarking modes

- CREATED_TIMESTAMP
- MODIFIED_TIMESTAMP




- This can be used for one of the ways introduced last section to ensure a file is new
 - Other options: Delete, move, filter
 - **Save and compare the file creation or modification times**

All contents © MuleSoft Inc.


17

17

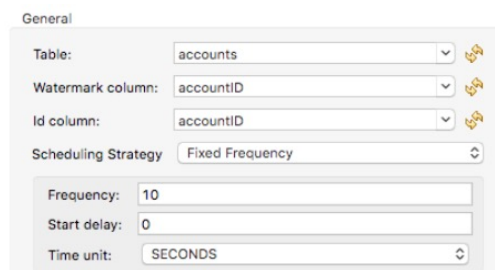
Triggering a flow for each row in a database table



- The Database connector has an **On Table Row** operation that is triggered for every row in a table



- The operation can handle
 - Generating the query
 - Watermarking
 - Idempotency across concurrent requests
- You can specify one, both, or neither of
 - Watermark column
 - Id column



All contents © MuleSoft Inc.

18


18

Using automatic watermarking with database tables



- When a watermark column is specified, this query is automatically generated and used

**SELECT * FROM table
WHERE table_column > :watermark**



On Table Row

- On each poll, the component will go through all the retrieved rows and store the maximum value obtained

General

Table:	accounts
Watermark column:	accountID
Id column:	accountID
Scheduling Strategy	Fixed Frequency
Frequency:	10
Start delay:	0
Time unit:	SECONDS

All contents © MuleSoft Inc.

19


19

Handling idempotency across concurrent requests



- A new poll can be executed before the watermark is updated if**

- The poll interval is small
- The amount of rows is big
- Processing one single row takes too much time



On Table Row

- To avoid a record being processed more than once**

- Specify an ID column
 - A unique identifier for the row
- The listener will make sure the row is not processed again if
 - It has already been retrieved and
 - Processing of it hasn't finished yet

General

Table:	accounts
Watermark column:	accountID
Id column:	accountID
Scheduling Strategy	Fixed Frequency
Frequency:	10
Start delay:	0
Time unit:	SECONDS

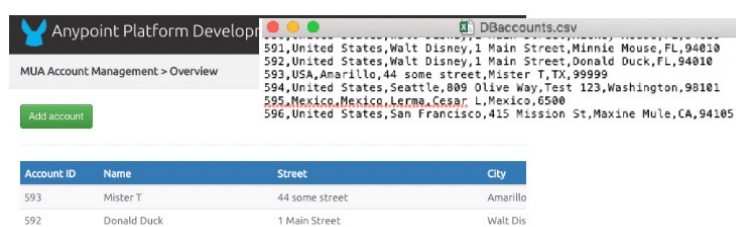
All contents © MuleSoft Inc.

20

20

Walkthrough 12-2: Trigger a flow when a new record is added to a database & use automatic watermarking

- Add and configure a Database listener to check a table on a set frequency for new records
- Use the listener's automatic watermarking to track the ID of the latest record retrieved and trigger the flow whenever a new record is added
- Output new records to a CSV file
- Use a form to add a new account to the table and see the CSV updated



The screenshot shows the Anypoint Platform Developer console. On the left, the 'MUA Account Management > Overview' page is visible, featuring an 'Add account' button. On the right, a CSV file named 'DBAccounts.csv' is displayed, containing the following data:

Account ID	Name	Street	City
593	Mister T	44 some street	Amarillo
592	Donald Duck	1 Main Street	Walt Dis

All contents © MuleSoft Inc.

21

21

Using manual watermarking and scheduling flows



22

Handling watermarking manually



- The general process
 - Schedule when a flow should be executed
 - Give the watermark a default value
 - On the first sync
 - Determine a new watermark value
 - Store the watermark value so it is available in the future to other flow executions
 - On later syncs
 - Retrieve the watermark from storage
 - Check if each item in the data set should be retrieved based on the watermark value

All contents © MuleSoft Inc.

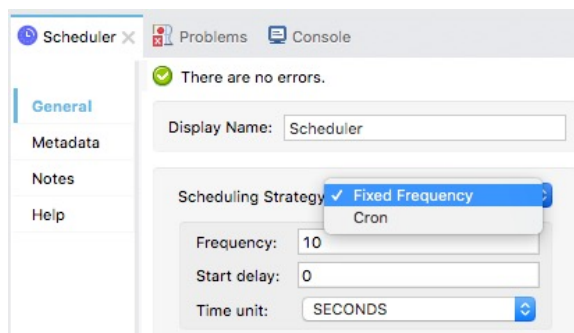

23

23

Triggering flows at a certain time or frequency



- Some connector event sources use a scheduling strategy to trigger a flow
 - Like On New or Updated File and On Table Row
- To trigger **any** flow at **any** time, use the **Scheduler** event source



All contents © MuleSoft Inc.

24

24

Two types of scheduling strategies



- **Fixed frequency**

- The default is to poll every 1000 milliseconds

- **Cron**

- A standard for describing time and date information
- Can specify either
 - An event to occur just once at a certain time
 - A recurring event on some frequency

0 15 10 ? * *

Poll at 10:15am every day

0 15 10 * * ? 2018

Poll at 10:15pm every day in 2018

1 1 1 1 1,6 ?


Poll the first day of January and June every year in the first second of the first minute of the first hour

All contents © MuleSoft Inc.

25

25

Persisting data across executions of flows



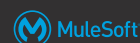
- Use the **Object Store** component to store simple key-value pairs
 - The component was designed to store
 - Synchronization information like watermarks
 - Temporal information like access tokens
 - User information
 - Retrieved values can be accessed through storage in a target variable
 - This storage causes the component to output the same message as the one received
- Each Mule application has an Object Store that is
 - Available without any setup or configuration
 - Persistent
 - Saved to file for embedded Mule and standalone Mule runtime
 - Saved to data storage for CloudHub
 - Saved to shared distributed memory for clustered Mule runtimes

All contents © MuleSoft Inc.

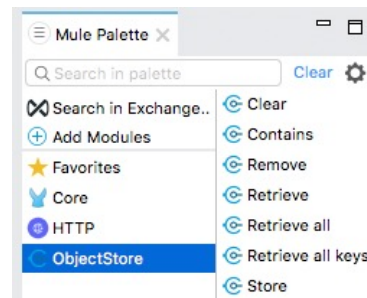
26

26

Using the Object Store connector for watermarking



- Add the ObjectStore module to the project
- Use the **Retrieve** operation to retrieve a watermark value and to assign a default value for the first poll
- Use the watermark value in a processor to retrieve the desired items
 - Like in a database query for records in a table
- Use the **Store** operation to determine and store a watermark value



All contents © MuleSoft Inc.

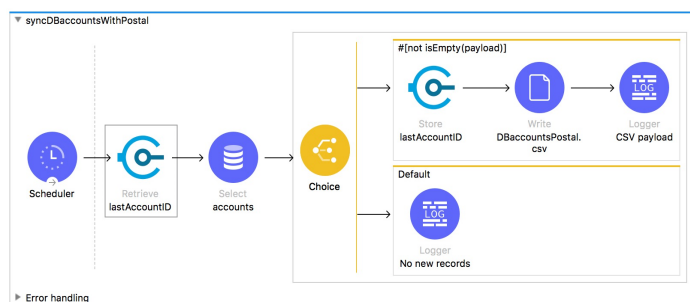
27

27

Walkthrough 12-3: Schedule a flow and use manual watermarking



- Use the Scheduler component to create a new flow that executes at a specific frequency
- Retrieve accounts with a specific postal code from the accounts table
- Use the Object Store component to store the ID of the latest record and then use it to only retrieve new records



All contents © MuleSoft Inc.

28

28

Publishing and consuming JMS messages



29

Java Messaging Service (JMS)



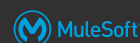
- Is a widely-used API for enabling applications to communicate through the exchange of messages
- Simplifies application development by providing a standard interface for creating, sending, and receiving messages
- Supports two messaging models
 - **Queues:** PTP (point-to-point or 1:1)
 - A sender sends messages to a queue & a single receiver pulls the message off the queue
 - The receiver does not need to be listening to the queue at the time the message is sent
 - **Topics:** Pub-Sub (publish/subscribe or 1:many)
 - A publisher sends a message to a topic & all active subscribers receive the message
 - Subscribers not actively listening will miss the published message (unless messages are made durable)

All contents © MuleSoft Inc.

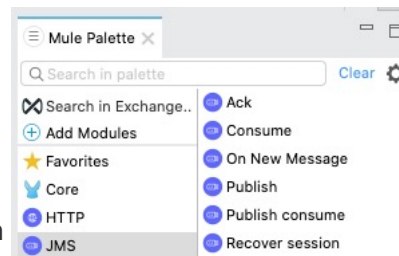
30

30

Using the JMS connector



- Add the JMS module to the project
- Configure a global element configuration
 - By default, it is set up with a finely tuned set of values for both for publishing and consuming messages
 - Typically, you just need to configure which connection should be used
- Use operations to publish and/or consume messages to destinations




All contents © MuleSoft Inc.

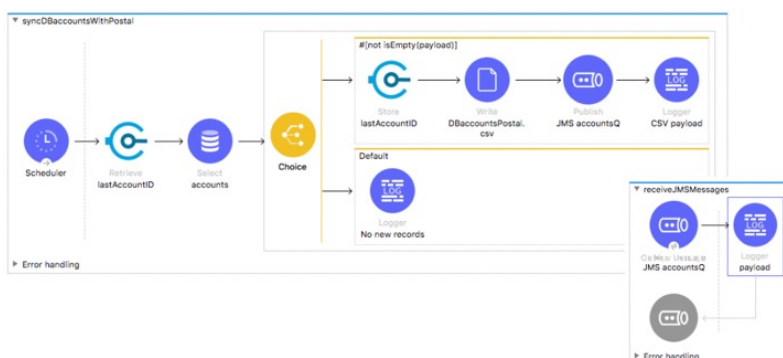
31

31

Walkthrough 12-4: Publish and listen for JMS messages



- Add and configure a JMS connector for ActiveMQ (that uses an in-memory broker)
- Send messages to a JMS queue
- Listen for and process messages from a JMS queue

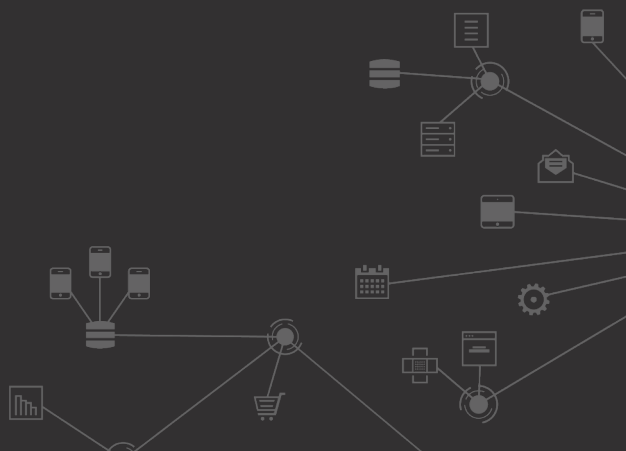


All contents © MuleSoft Inc.

32

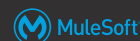
32

Summary



33

Summary



- Use **watermarks** to synchronize data across data stores
 - Use either **manual** or the **automatic** watermarking available for some connectors
- Use the family of **File**, **FTP**, **FTPS**, and **SFTP** connectors to work with files and folders
- Use the **On New or Updated File** listener to trigger flows when files are added, created, or updated
 - Use the connector's **automatic watermarking** to determine if a file is new based on a creation or modification timestamp
- Use the **On Table Row** listener when new records are added to a database table
 - Use the connector's **automatic watermarking** to determine if the record is new

All contents © MuleSoft Inc.

34

34

Summary



- Use the **Scheduler** component to schedule flows to run at a certain time or frequency
 - Use a watermark to keep a persistent variable between scheduling events
- Use the **Object Store** connector to persist and share a watermark (or other data) across flow executions
- Use the **JMS** connector to publish and consume messages
 - Connect to any JMS messaging service that implements the JMS spec